

Lab 3.5 - Student Notebook

Importing the data

By running the following cells, the data will be imported and ready for use.

Note: The following cells represent the key steps in the previous labs.

```
In [1]: bucket='c169682a4380821111163942t1w260663698571-labbucket-8hkf2hcmkeqb'
```

```
In [2]: import warnings, requests, zipfile, io
warnings.simplefilter('ignore')
import pandas as pd
from scipy.io import arff

import os
import boto3
import sagemaker
from sagemaker.image_uris import retrieve
from sklearn.model_selection import train_test_split
```

```
sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /home/ec2-user/.config/sagemaker/config.yaml
```

```
In [3]: f_zip = 'http://archive.ics.uci.edu/ml/machine-learning-databases/00212/vertebra
r = requests.get(f_zip, stream=True)
Vertebral_zip = zipfile.ZipFile(io.BytesIO(r.content))
Vertebral_zip.extractall()

data = arff.loadarff('column_2C_weka.arff')
df = pd.DataFrame(data[0])

class_mapper = {b'Abnormal':1,b'Normal':0}
df['class']=df['class'].replace(class_mapper)

cols = df.columns.tolist()
cols = cols[-1:] + cols[:-1]
df = df[cols]

train, test_and_validate = train_test_split(df, test_size=0.2, random_state=42,
test, validate = train_test_split(test_and_validate, test_size=0.5, random_state

prefix='lab3'

train_file='vertebral_train.csv'
test_file='vertebral_test.csv'
validate_file='vertebral_validate.csv'

s3_resource = boto3.Session().resource('s3')
def upload_s3_csv(filename, folder, dataframe):
```

```

csv_buffer = io.StringIO()
dataframe.to_csv(csv_buffer, header=False, index=False )
s3_resource.Bucket(bucket).Object(os.path.join(prefix, folder, filename)).put(
    csv_buffer.getvalue()

upload_s3_csv(train_file, 'train', train)
upload_s3_csv(test_file, 'test', test)
upload_s3_csv(validate_file, 'validate', validate)

container = retrieve('xgboost', boto3.Session().region_name, '1.0-1')

hyperparams={"num_round": "42",
              "eval_metric": "auc",
              "objective": "binary:logistic"}

s3_output_location="s3://{}/{}/output/".format(bucket,prefix)
xgb_model=sagemaker.estimator.Estimator(container,
                                         sagemaker.get_execution_role(),
                                         instance_count=1,
                                         instance_type='ml.m4.xlarge',
                                         output_path=s3_output_location,
                                         hyperparameters=hyperparams,
                                         sagemaker_session=sagemaker.Session())

train_channel = sagemaker.inputs.TrainingInput(
    "s3://{}/{}/train/".format(bucket,prefix,train_file),
    content_type='text/csv')

validate_channel = sagemaker.inputs.TrainingInput(
    "s3://{}/{}/validate/".format(bucket,prefix,validate_file),
    content_type='text/csv')

data_channels = {'train': train_channel, 'validation': validate_channel}

xgb_model.fit(inputs=data_channels, logs=False)

print('ready for hosting!')

```

```

INFO:sagemaker:Creating training-job with name: sagemaker-xgboost-2025-08-10-07-34-27-290

```

```

2025-08-10 07:34:28 Starting - Starting the training job..
2025-08-10 07:34:43 Starting - Preparing the instances for training...
2025-08-10 07:35:07 Downloading - Downloading input data.....
2025-08-10 07:35:37 Downloading - Downloading the training image.....
2025-08-10 07:36:33 Training - Training image download completed. Training in progress.....
2025-08-10 07:36:54 Uploading - Uploading generated training model.
2025-08-10 07:37:07 Completed - Training job completed
ready for hosting!

```

Step 1: Hosting the model

Now that you have a trained model, you can host it by using Amazon SageMaker hosting services.

The first step is to deploy the model. Because you have a model object, *xgb_model*, you can use the **deploy** method. For this lab, you will use a single ml.m4.xlarge instance.

```
In [4]: xgb_predictor = xgb_model.deploy(initial_instance_count=1,
      serializer = sagemaker.serializers.CSVSerializer(),
      instance_type='ml.m4.xlarge')
```

```
INFO:sagemaker:Creating model with name: sagemaker-xgboost-2025-08-10-07-37-09-124
INFO:sagemaker:Creating endpoint-config with name sagemaker-xgboost-2025-08-10-07-37-09-124
INFO:sagemaker:Creating endpoint with name sagemaker-xgboost-2025-08-10-07-37-09-124
-----!
```

Step 2: Performing predictions

Now that you have a deployed model, you will run some predictions.

First, review the test data and re-familiarize yourself with it.

```
In [5]: test.shape
```

```
Out[5]: (31, 7)
```

You have 31 instances, with seven attributes. The first five instances are:

```
In [6]: test.head(5)
```

```
Out[6]:
```

	class	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	degree
136	1	88.024499	39.844669	81.774473	48.179830	116.601538	
230	0	65.611802	23.137919	62.582179	42.473883	124.128001	
134	1	52.204693	17.212673	78.094969	34.992020	136.972517	
130	1	50.066786	9.120340	32.168463	40.946446	99.712453	
47	1	41.352504	16.577364	30.706191	24.775141	113.266675	

You don't need to include the target value (class). This predictor can take data in the comma-separated values (CSV) format. You can thus get the first row *without the class column* by using the following code:

```
test.iloc[:1,1:]
```

The **iloc** function takes parameters of *[rows,cols]*

To only get the first row, use `0:1` . If you want to get row 2, you could use `1:2` .

To get all columns *except* the first column (col 0), use `1:`

```
In [7]: row = test.iloc[0:1,1:]
      row.head()
```

Out[7]:

	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	degree_spoi
136	88.024499	39.844669	81.774473	48.17983	116.601538	

You can convert this to a comma-separated values (CSV) file, and store it in a string buffer.

```
In [8]: batch_X_csv_buffer = io.StringIO()
row.to_csv(batch_X_csv_buffer, header=False, index=False)
test_row = batch_X_csv_buffer.getvalue()
print(test_row)
```

88.0244989,39.84466878,81.77447308,48.17983012,116.6015376,56.76608323

Now, you can use the data to perform a prediction.

```
In [9]: xgb_predictor.predict(test_row)
```

Out[9]: b'0.9966071844100952'

The result you get isn't a 0 or a 1. Instead, you get a *probability score*. You can apply some conditional logic to the probability score to determine if the answer should be presented as a 0 or a 1. You will work with this process when you do batch predictions.

For now, compare the result with the test data.

```
In [10]: test.head(5)
```

Out[10]:

	class	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	degre
136	1	88.024499	39.844669	81.774473	48.179830	116.601538	
230	0	65.611802	23.137919	62.582179	42.473883	124.128001	
134	1	52.204693	17.212673	78.094969	34.992020	136.972517	
130	1	50.066786	9.120340	32.168463	40.946446	99.712453	
47	1	41.352504	16.577364	30.706191	24.775141	113.266675	

Question: Is the prediction accurate?

Challenge task: Update the previous code to send the second row of the dataset. Are those predictions correct? Try this task with a few other rows.

It can be tedious to send these rows one at a time. You could write a function to submit these values in a batch, but SageMaker already has a batch capability. You will examine that feature next. However, before you do, you will terminate the model.

Step 3: Terminating the deployed model

To delete the endpoint, use the **delete_endpoint** function on the predictor.

```
In [11]: xgb_predictor.delete_endpoint(delete_endpoint_config=True)
```

```
INFO:sagemaker:Deleting endpoint configuration with name: sagemaker-xgboost-2025-08-10-07-37-09-124
```

```
INFO:sagemaker:Deleting endpoint with name: sagemaker-xgboost-2025-08-10-07-37-09-124
```

Step 4: Performing a batch transform

When you are in the training-testing-feature engineering cycle, you want to test your holdout or test sets against the model. You can then use those results to calculate metrics. You could deploy an endpoint as you did earlier, but then you must remember to delete the endpoint. However, there is a more efficient way.

You can use the transformer method of the model to get a transformer object. You can then use the transform method of this object to perform a prediction on the entire test dataset. SageMaker will:

- Spin up an instance with the model
- Perform a prediction on all the input values
- Write those values to Amazon Simple Storage Service (Amazon S3)
- Finally, terminate the instance

You will start by turning your data into a CSV file that the transformer object can take as input. This time, you will use **iloc** to get all the rows, and all columns *except* the first column.

```
In [12]: batch_X = test.iloc[:,1:];
batch_X.head()
```

```
Out[12]:
```

	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	degree_spoi
136	88.024499	39.844669	81.774473	48.179830	116.601538	
230	65.611802	23.137919	62.582179	42.473883	124.128001	
134	52.204693	17.212673	78.094969	34.992020	136.972517	
130	50.066786	9.120340	32.168463	40.946446	99.712453	
47	41.352504	16.577364	30.706191	24.775141	113.266675	

Next, write your data to a CSV file.

```
In [13]: batch_X_file='batch-in.csv'
upload_s3_csv(batch_X_file, 'batch-in', batch_X)
```

Last, before you perform a transform, configure your transformer with the input file, output location, and instance type.

```
In [14]: batch_output = "s3://{}/{}batch-out/".format(bucket,prefix)
batch_input = "s3://{}/{}batch-in/{}".format(bucket,prefix,batch_X_file)

xgb_transformer = xgb_model.transformer(instance_count=1,
                                         instance_type='ml.m4.xlarge',
                                         strategy='MultiRecord',
                                         assemble_with='Line',
                                         output_path=batch_output)

xgb_transformer.transform(data=batch_input,
                          data_type='S3Prefix',
                          content_type='text/csv',
                          split_type='Line')

xgb_transformer.wait()
```

```
INFO:sagemaker:Creating model with name: sagemaker-xgboost-2025-08-10-07-40-41-617
INFO:sagemaker:Creating transform job with name: sagemaker-xgboost-2025-08-10-07-40-42-197
```

```
.....
...
```

After the transform completes, you can download the results from Amazon S3 and compare them with the input.

First, download the output from Amazon S3 and load it into a pandas DataFrame.

```
In [15]: s3 = boto3.client('s3')
obj = s3.get_object(Bucket=bucket, Key="{}/batch-out/{}".format(prefix,'batch-in'))
target_predicted = pd.read_csv(io.BytesIO(obj['Body'].read()),sep=',',names=['class'])
target_predicted.head(5)
```

```
Out[15]:
```

	class
0	0.996607
1	0.777283
2	0.994641
3	0.993690
4	0.939139

You can use a function to convert the probability into either a 0 or a 1.

The first table output will be the *predicted values*, and the second table output is the *original test data*.

```
In [16]: def binary_convert(x):
threshold = 0.65
if x > threshold:
    return 1
else:
    return 0

target_predicted['binary'] = target_predicted['class'].apply(binary_convert)
```

```
print(target_predicted.head(10))
test.head(10)
```

```
      class  binary
0  0.996607      1
1  0.777283      1
2  0.994641      1
3  0.993690      1
4  0.939139      1
5  0.997396      1
6  0.991977      1
7  0.987518      1
8  0.993334      1
9  0.682776      1
```

Out[16]:

	class	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	degre
--	-------	------------------	-------------	-----------------------	--------------	---------------	-------

136	1	88.024499	39.844669	81.774473	48.179830	116.601538
230	0	65.611802	23.137919	62.582179	42.473883	124.128001
134	1	52.204693	17.212673	78.094969	34.992020	136.972517
130	1	50.066786	9.120340	32.168463	40.946446	99.712453
47	1	41.352504	16.577364	30.706191	24.775141	113.266675
135	1	77.121344	30.349874	77.481083	46.771470	110.611148
100	1	84.585607	30.361685	65.479486	54.223922	108.010218
89	1	71.186811	23.896201	43.696665	47.290610	119.864938
297	0	45.575482	18.759135	33.774143	26.816347	116.797007
4	1	49.712859	9.652075	28.317406	40.060784	108.168725

Note: The *threshold* in the **binary_convert** function is set to .65.

Challenge task: Experiment with changing the value of the threshold. Does it impact the results?

Note: The initial model might not be good. You will generate some metrics in the next lab, before you tune the model in the final lab.

Congratulations!

You have completed this lab, and you can now end the lab by following the lab guide instructions.

CHALLENGE

```
In [17]: def binary_convert(x):
          threshold = 0.89
          if x > threshold:
              return 1
          else:
              return 0
```

```
target_predicted['binary'] = target_predicted['class'].apply(binary_convert)

print(target_predicted.head(10))
test.head(10)
```

	class	binary
0	0.996607	1
1	0.777283	0
2	0.994641	1
3	0.993690	1
4	0.939139	1
5	0.997396	1
6	0.991977	1
7	0.987518	1
8	0.993334	1
9	0.682776	0

Out[17]:

	class	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	deg
--	-------	------------------	-------------	-----------------------	--------------	---------------	-----

136	1	88.024499	39.844669	81.774473	48.179830	116.601538	
230	0	65.611802	23.137919	62.582179	42.473883	124.128001	
134	1	52.204693	17.212673	78.094969	34.992020	136.972517	
130	1	50.066786	9.120340	32.168463	40.946446	99.712453	
47	1	41.352504	16.577364	30.706191	24.775141	113.266675	
135	1	77.121344	30.349874	77.481083	46.771470	110.611148	
100	1	84.585607	30.361685	65.479486	54.223922	108.010218	
89	1	71.186811	23.896201	43.696665	47.290610	119.864938	
297	0	45.575482	18.759135	33.774143	26.816347	116.797007	
4	1	49.712859	9.652075	28.317406	40.060784	108.168725	

```
In [18]: def binary_convert(x):
          threshold = 0.9
          if x > threshold:
              return 1
          else:
              return 0

          target_predicted['binary'] = target_predicted['class'].apply(binary_convert)

          print(target_predicted.head(10))
          test.head(10)
```

	class	binary
0	0.996607	1
1	0.777283	0
2	0.994641	1
3	0.993690	1
4	0.939139	1
5	0.997396	1
6	0.991977	1
7	0.987518	1
8	0.993334	1
9	0.682776	0

Out[18]:

	class	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	deg
136	1	88.024499	39.844669	81.774473	48.179830	116.601538	
230	0	65.611802	23.137919	62.582179	42.473883	124.128001	
134	1	52.204693	17.212673	78.094969	34.992020	136.972517	
130	1	50.066786	9.120340	32.168463	40.946446	99.712453	
47	1	41.352504	16.577364	30.706191	24.775141	113.266675	
135	1	77.121344	30.349874	77.481083	46.771470	110.611148	
100	1	84.585607	30.361685	65.479486	54.223922	108.010218	
89	1	71.186811	23.896201	43.696665	47.290610	119.864938	
297	0	45.575482	18.759135	33.774143	26.816347	116.797007	
4	1	49.712859	9.652075	28.317406	40.060784	108.168725	

In [19]:

```
def binary_convert(x):
    threshold = 0.99
    if x > threshold:
        return 1
    else:
        return 0

target_predicted['binary'] = target_predicted['class'].apply(binary_convert)

print(target_predicted.head(10))
test.head(10)
```

	class	binary
0	0.996607	1
1	0.777283	0
2	0.994641	1
3	0.993690	1
4	0.939139	0
5	0.997396	1
6	0.991977	1
7	0.987518	0
8	0.993334	1
9	0.682776	0

Out[19]:

	class	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	deg
136	1	88.024499	39.844669	81.774473	48.179830	116.601538	
230	0	65.611802	23.137919	62.582179	42.473883	124.128001	
134	1	52.204693	17.212673	78.094969	34.992020	136.972517	
130	1	50.066786	9.120340	32.168463	40.946446	99.712453	
47	1	41.352504	16.577364	30.706191	24.775141	113.266675	
135	1	77.121344	30.349874	77.481083	46.771470	110.611148	
100	1	84.585607	30.361685	65.479486	54.223922	108.010218	
89	1	71.186811	23.896201	43.696665	47.290610	119.864938	
297	0	45.575482	18.759135	33.774143	26.816347	116.797007	
4	1	49.712859	9.652075	28.317406	40.060784	108.168725	

In [20]:

```
def binary_convert(x):
    threshold = 0.999
    if x > threshold:
        return 1
    else:
        return 0

target_predicted['binary'] = target_predicted['class'].apply(binary_convert)

print(target_predicted.head(10))
test.head(10)
```

	class	binary
0	0.996607	0
1	0.777283	0
2	0.994641	0
3	0.993690	0
4	0.939139	0
5	0.997396	0
6	0.991977	0
7	0.987518	0
8	0.993334	0
9	0.682776	0

Out[20]:

	class	pelvic_incidence	pelvic_tilt	lumbar_lordosis_angle	sacral_slope	pelvic_radius	deg
136	1	88.024499	39.844669	81.774473	48.179830	116.601538	
230	0	65.611802	23.137919	62.582179	42.473883	124.128001	
134	1	52.204693	17.212673	78.094969	34.992020	136.972517	
130	1	50.066786	9.120340	32.168463	40.946446	99.712453	
47	1	41.352504	16.577364	30.706191	24.775141	113.266675	
135	1	77.121344	30.349874	77.481083	46.771470	110.611148	
100	1	84.585607	30.361685	65.479486	54.223922	108.010218	
89	1	71.186811	23.896201	43.696665	47.290610	119.864938	
297	0	45.575482	18.759135	33.774143	26.816347	116.797007	
4	1	49.712859	9.652075	28.317406	40.060784	108.168725	

In []:

